

# Change Point Detection in Genomic Data

Alan Chau   Ana Ignatieva  
James Thornton   Lorenzo Pacchiardi

02/11/2018

# Long Chain DNA Sequences




Figure 1: Plot of a DNA with 5000 base pair

- Seemingly long and messy spectrum of A, T, C and G. Can you observe any structure?

# Long Chain DNA Sequences




Figure 2: Plot of a DNA with 5000 base pair

- I found one after staring at the data for 2 days.

# Long Chain DNA Sequences




Figure 2: Plot of a DNA with 5000 base pair

- I found one after staring at the data for 2 days.
- Is there a systematic way of doing this structure discovery?

# Long Chain DNA Sequences




Figure 2: Plot of a DNA with 5000 base pair

- I found one after staring at the data for 2 days.
- Is there a systematic way of doing this structure discovery?
- In other words, can we detect change points systematically along this sequence?

# Talk Overview

- Alan Chau - Exact algorithms for detections
- Lorenzo Pacchiardi - Maximum Likelihood and approximate search methods
- Ana Ignatieva - Hidden Markov Model
- James Thornton - Bayesian Model and Grid-Merge

# Problem Statement

- Consider a sequence of data  $\mathbf{y} = \{y_t\}_{t=1}^T$  where  $y_t$  can be continuous, discrete or categorical.
- Configure the set of possible change points as  $\mathcal{T}_T := \{\tau : 0 = \tau_0 < \tau_1 < \dots < \tau_K < \tau_{K+1} = T\}$
- We aim to minimise the following:

$$V(\tau, \mathbf{y}) = \min_{\tau \in \mathcal{T}_T} \sum_{k=0}^K \mathcal{C}(y_{\tau_k+1:\tau_{k+1}}) + \text{Pen}\{\tau, \beta\} \quad (1)$$

where  $\mathcal{C}$  is some cost function measuring heterogeneity of the segment.

# Optimal Partitioning

- One obvious choice is to pick the  $L_0$  norm,  $\beta K$ .
- Set  $F(T)$  as  $V(\tau, \mathbf{y})$  restricted to the domain  $\mathcal{T}_T$ , we then realise:

$$F(T) = \min_{\tau \in \mathcal{T}_T} \left\{ \sum_{k=0}^K [\mathcal{C}(y_{\tau_k+1:\tau_{k+1}}) + \beta] \right\} \quad (2)$$

$$= \min_t \left\{ \sum_{k=0}^{K-1} [\mathcal{C}(y_{\tau_k+1:\tau_{k+1}}) + \beta] + \mathcal{C}(y_{t+1:T}) + \beta \right\} \quad (3)$$

$$= \min_t \{F(t) + \mathcal{C}(y_{t+1:T}) + \beta\} \quad (4)$$

- We can then use a dynamic program and start from  $F(1)$  and obtain an exact solution in  $\mathcal{O}(T^2)$



# Pruned Exact Linear Time (PELT)

- We get exactness but it is slow.

# Pruned Exact Linear Time (PELT)

- We get exactness but it is slow.
- We can use pruning to reduce complexity and achieve approximate linear run time.

# Pruned Exact Linear Time (PELT)

- We get exactness but it is slow.
- We can use pruning to reduce complexity and achieve approximate linear run time.

## Theorem

*We assume there exists a constant  $K$  such that for all  $t < s < T$ ,*

$$\mathcal{C}(y_{t+1:s}) + \mathcal{C}(y_{s+1:T}) + K \leq \mathcal{C}(y_{t+1:T})$$

*Then if*

$$F(t) + \mathcal{C}(y_{t+1:s}) + K \geq F(s)$$

*holds, at a future time  $T > s$ ,  $t$  can never be the optimal last change point prior to  $T$ .*

# Pruned Exact Linear Time (PELT)

## Theorem

*We assume there exists a constant  $K$  such that for all  $t < s < T$ ,*

$$\mathcal{C}(y_{t+1:s}) + \mathcal{C}(y_{s+1:T}) + K \leq \mathcal{C}(y_{t+1:T})$$

*Then if*

$$F(t) + \mathcal{C}(y_{t+1:s}) + K \geq F(s)$$

*holds, at a future time  $T > s$ ,  $t$  can never be the optimal last change point prior to  $T$ .*

- Intuitively speaking, if the condition is met, then  $s$  is always a better changepoint than  $t$ , thus we don't have to consider that anymore in future steps.

# Pruned Exact Linear Time

**Initialise:** Let  $T$  be the number of data, set  $F(0) = -\beta$ ,  $cp(0) = NULL$  and  $R_1 = \{0\}$

**Iterate:** For  $\tau^* = 1, \dots, T$

- ① Compute  $F(\tau^*) = \min_{\tau \in R_{\tau^*}} [F(\tau) + \mathcal{C}(y_{\tau+1} : \tau^*) + \beta]$
- ② Let  $\tau' = \arg \min_{\tau \in R_{\tau^*}} [F(\tau) + \mathcal{C}(y_{\tau+1} : \tau^*) + \beta]$
- ③ set  $cp(\tau^*) = [cp(\tau'), \tau']$
- ④ set  $R_{\tau^*+1} = \{\tau \in R_{\tau^*} \cap \tau^* : F(\tau) + \mathcal{C}(y_{\tau+1} : \tau^*) + K \leq F(\tau^*)\}$

**Output:** A vector of change points in  $cp(T)$

**Algorithm 1:** PELT Algorithm

# Application


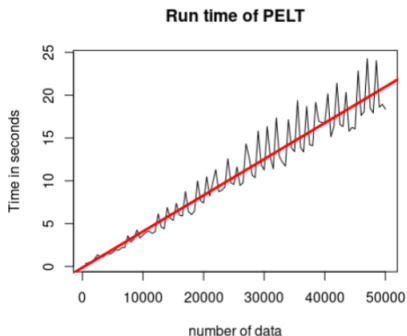



Figure 3: PELT applied to a sequence of DNA with length 200

# Application



**Figure 4:** (left) PELT applied to the first 200 DNA sequence. (right) Run time analysis on PELT

# Problem with large sequences




Figure 5: Limitation on the PELT method with large sequence



# Pruned Exact Linear Time with Modified $L_0$ penalty

To tackle this concentration of change points, we introduce a modified penalty for the PELT algorithm to control the spread,

$$\text{Pen}_{mL_0}(\tau, \beta) = 3K\beta \log(T) + \beta \log(\beta) \sum_{k=0}^K \log\left(\frac{\tau_{k+1} - \tau_k}{T}\right) \quad (5)$$

This can be easily incorporate to PELT by setting:

$$\mathcal{C}'(y_{a+1:b}) \leftarrow \mathcal{C}(y_{a+1:b}) + \beta \log(\beta) \log\left(\frac{b-a}{T}\right) \quad \beta' \leftarrow 3\beta \log(T) \quad (6)$$

# Results




Figure 6: More evenly spread change points

# MLE of transition matrix

Consider a standard Markov Chain. MLE for the transition probability:


$$\hat{p}_{ij} = \frac{n_{ij}}{\sum_{j=1}^m n_{ij}} = \frac{n_{ij}}{n_i} \longrightarrow p_{ij}$$

$n_{ij}$ : number of observed transitions  $i \rightarrow j$

# MLE of transition matrix

If the inferred segment is actually the composition of two MCs with different transition probabilities, we get:

$$\hat{p}_{ij} \longrightarrow \frac{N_0 p_i^0 p_{ij}^0 + N_1 p_i^1 p_{ij}^1}{N_0 p_i^0 + N_1 p_i^1}, \quad N_0, N_1 \rightarrow \infty$$



# MLE of transition matrix

If the inferred segment is actually the composition of two MCs with different transition probabilities, we get:


$$\hat{p}_{ij} \longrightarrow \frac{N_0 p_i^0 p_{ij}^0 + N_1 p_i^1 p_{ij}^1}{N_0 p_i^0 + N_1 p_i^1}, \quad N_0, N_1 \rightarrow \infty$$



$\hat{p}_{ij} \rightarrow p_{ij}^0 \iff N_1/N_0 \rightarrow 0$ . This is usually not guaranteed in change point detection algorithms.



# Markov model of genome segmentation (Thakur et al., 2007)

- Recursive binary segmentation procedure.
- Maximizes Jensen Shannon divergence between the two subsequences  $S_1, S_2$ :  $D_{JS}(S_1, S_2) = H(S) - \pi_1 H(S_1) - \pi_2 H(S_2)$
- Each segmentation is accepted if it satisfies the BIC criterion  $\Delta \mathcal{C}_{BIC} < 0 \iff 2ND_{JS}(S_1, S_2) > 16 \ln N$




# Experiment on the genomic data

- Time complexity  $\mathcal{O}(N \log_2 K)$ ,  $K$ : number of change points.
- Algorithm was run on the first 5 million nucleotides. Segments for the first 1500 are reported in picture.



## HMM



- $K$  hidden states, transition matrix  $A$
- Emissions  $y_1, \dots, y_T$ , transitions  $P_{ij}^{(s)}$
- Changepoints = hidden state transitions
- Used for genome data (Churchill, 1989; Peshkin and S. Gelfand, 1999)



# Hidden state transitions

For changepoints want structure of  $A$ :

$$A = \begin{pmatrix} 1 - \lambda_1 & \lambda_1 & 0 & \dots & 0 \\ 0 & 1 - \lambda_2 & \lambda_2 & \dots & 0 \\ \vdots & & \ddots & \ddots & 0 \\ 0 & \dots & 0 & 1 - \lambda_{K-1} & \lambda_{K-1} \\ 0 & \dots & 0 & 0 & 1 \end{pmatrix}$$

- Small probability of transitioning to next state
- Can't revisit previous states
- Can't jump ahead

# Parameter fitting


Baum-Welch (Durbin et al., 1998):

- Version of E-M
- Need to incorporate Markov dependency of emissions
- Iteratively compute  $\mathbb{P}(S_t|y_1, \dots, y_T)$
- Expected number of times  $(ij)$  appears in state  $s$   
 $\implies$  update  $P_{ij}^{(s)}$
- Expected number of state transitions  $i \rightarrow j$   
 $\implies$  update  $A_{ij}$

Decode the HMM:

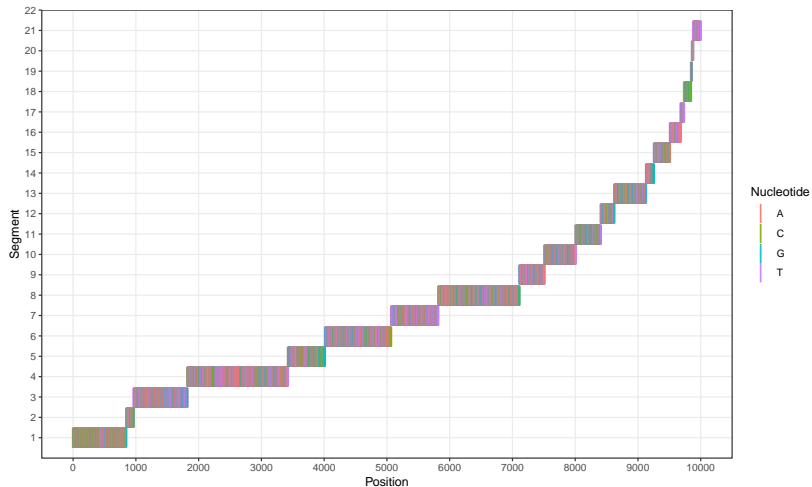
- Viterbi: get most likely path through states
- Posterior  $\mathbb{P}(S_t|y_1, \dots, y_T)$

# Simulated data



- Actual changepoints at every 1,000
- Almost!

# Real data

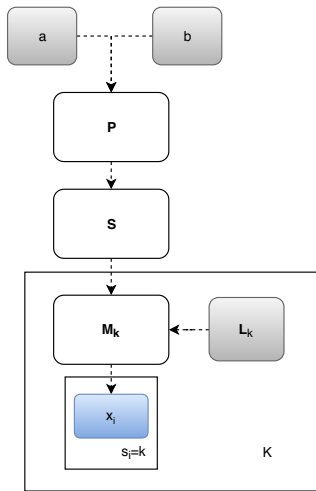


- Initial guess for max number of changepoints = 50

# Conclusions

- Flexible:
  - number of changepoints
  - incorporate information into structure of  $A$
- General class of models
- Relatively slow

# Bayesian Change Point Detection



Generative Model:

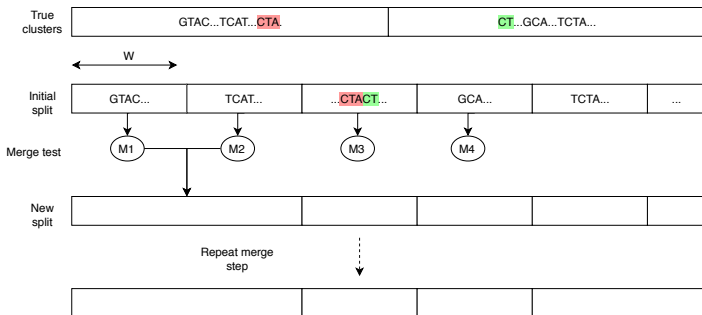
- $P \sim \text{BetaDiag}(a, b)$
- $S \sim \text{MarkovChain}(P)$  where  $s_1 = 1$
- $M_k | S \sim \text{DirMat}(\lambda) \forall k \in [1 : K]$
- $x_{1:n_k}^k | S, M_k \sim \text{MarkovChain}(M_k)$

Implementation:

- Conjugate
- MCMC
  - Blocked Gibbs Sampling
  - Simulated annealing

# Grid-Merge Heuristic


Split chain index  $\rightarrow$  Merge on distance between MLE transition matrices  
 $\rightarrow \dots$




- Applicable for large data-sets
- Intuitive interpretation
- Approximation error

# Results on Simulated Data


Bayesian Model



Grid-Merge



True Change Points





- Gary A Churchill. Stochastic models for heterogeneous DNA sequences. *Bulletin of mathematical biology*, 51(1):79–94, 1989.
- Richard Durbin, Sean R Eddy, Anders Krogh, and Graeme Mitchison. *Biological sequence analysis: probabilistic models of proteins and nucleic acids*. Cambridge university press, 1998.
- Leonid Peshkin and Mikhail S. Gelfand. Segmentation of yeast DNA using hidden Markov models. *Bioinformatics*, 15(12):980–986, 1999.
- Vivek Thakur, Rajeev K Azad, and Ram Ramaswamy. Markov models of genome segmentation. *Physical Review E*, 75(1):011915, 2007.